

# Querying with SQLAlchemy

---



**Xavier Morera**

HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA

@xmorera [www.xaviermorera.com](http://www.xaviermorera.com)



# Querying with SQLAlchemy




Picking a database



Connectors and connection strings



Querying using SQL vs. ORM



Joins, and hierarchical tables



# Picking a Database: Which One?

SQLite

PostgreSQL

MySQL

Oracle

Microsoft  
SQL Server

Or another one...





**Powerful open source database**

**Runs on virtually all major platforms**

- Top 3 of most widely used

**Client-server model**

**MariaDB too**





## Database engine

- Store and work with relational data

## No need for a database server

- Simple to use and portable

## Does not scale well

PostgreSQL



**Object-relational db management system**

- Advanced and open source

**Highly programmable and extensible**

**Extremely efficient concurrency support**

**Not as popular as MySQL**






Multi-model database management system

Commercial

- Widely used
- Enterprise





## Microsoft SQL Server



Widespread in the enterprise world

- On-premises and in the cloud

Efficient and easy to use










# Connecting to Databases with SQLAlchemy



Connector



Connection String

# Connecting to Databases with SQLAlchemy




Connector



Connection String

# Connecting to Databases with SQLAlchemy



Connectors



Connection String

# Connectors

SQLite

PostgreSQL

MySQL

Oracle

MS SQL Server

Or another one...



# Connectors

**SQLite**

pysqlite (sqlite3)

**PostgreSQL**

psycopg2

**MySQL**

MySQL Connector  
Python

**Oracle**

cx\_Oracle

**MS SQL Server**

PyODBC

More?



## SQLAlchemy 1.3 Documentation

Release: **1.3.0b2** **BETA RELEASE** | Release Date: January 25, 2019SQLAlchemy 1.3  
Documentation**BETA RELEASE**[Contents](#) | [Index](#)Search terms: 

ads via Carbon

## Dialects

[PostgreSQL](#)[MySQL](#)**SQLite**

- [Support for the SQLite database.](#)
- [Date and Time Types](#)
  - [Ensuring Text affinity](#)
- [SQLite Auto Incrementing Behavior](#)
  - [Using the AUTOINCREMENT Keyword](#)
  - [Allowing autoincrement behavior SQLAlchemy types other than Integer/INTEGER](#)
- [Database Locking Behavior /](#)

## SQLite

Support for the SQLite database.

## DBAPI Support

The following dialect/DBAPI options are available. Please refer to individual DBAPI sections for connect information.

- [pysqlite](#)
- [pysqlcipher](#)

## Date and Time Types

SQLite does not have built-in DATE, TIME, or DATETIME types, and pysqlite does not provide out of the box functionality for translating values between Python *datetime* objects and a SQLite-supported format. SQLAlchemy's own `DateTime` and related types provide date formatting and parsing functionality when SQLite is used. The implementation classes are `DATETIME`, `DATE` and `TIME`. These types represent dates and times as ISO formatted strings, which also nicely support ordering. There's no reliance on typical "libc" internals for these functions so historical dates are fully supported.

### Ensuring Text affinity

The DDL rendered for these types is the standard `DATE`, `TIME` and `DATETIME` indicators. However, custom storage formats can also be applied to these

## SQLAlchemy 1.3 Documentation

Release: **1.3.0b2** **BETA RELEASE** | Release Date: January 25, 2019SQLAlchemy 1.3  
Documentation**BETA RELEASE**[Contents](#) | [Index](#)Search terms: 

ads via Carbon

## Dialects

## PostgreSQL

- Support for the PostgreSQL database.
- Sequences/SERIAL/IDENTITY
  - PostgreSQL 10 IDENTITY columns
- Transaction Isolation Level
- Remote-Schema Table Introspection and PostgreSQL search\_path
- INSERT/UPDATE...RETURNING
- INSERT...ON CONFLICT (Upsert)
- Full Text Search
- FROM ONLY ...

## PostgreSQL

Support for the PostgreSQL database.

## DBAPI Support

The following dialect/DBAPI options are available. Please refer to individual DBAPI sections for connect information.

- [psycopg2](#)
- [pg8000](#)
- [psycopg2cffi](#)
- [py-postgresql](#)
- [pygresql](#)
- [zxJDBC](#) for Jython

## Sequences/SERIAL/IDENTITY

PostgreSQL supports sequences, and SQLAlchemy uses these as the default means of creating new primary key values for integer-based primary key columns. When creating tables, SQLAlchemy will issue the `SERIAL` datatype for integer-based primary key columns, which generates a sequence and server side default corresponding to the column.



## SQLAlchemy 1.3 Documentation

Release: **1.3.0b2** **BETA RELEASE** | Release Date: January 25, 2019SQLAlchemy 1.3  
Documentation**BETA RELEASE**[Contents](#) | [Index](#)Search terms: 

ads via Carbon

## Dialects

## PostgreSQL

**MySQL**

- [Support for the MySQL database.](#)
- [Supported Versions and Features](#)
- [Connection Timeouts and Disconnects](#)
- [CREATE TABLE arguments including Storage Engines](#)
- [Case Sensitivity and Table Reflection](#)
- [Transaction Isolation Level](#)
- [AUTO\\_INCREMENT Behavior](#)

## MySQL

Support for the MySQL database.

## DBAPI Support

The following dialect/DBAPI options are available. Please refer to individual DBAPI sections for connect information.

- [mysqlclient](#) (maintained fork of MySQL-Python)
- [PyMySQL](#)
- [MySQL Connector/Python](#)
- [CyMySQL](#)
- [OurSQL](#)
- [Google Cloud SQL](#)
- [PyODBC](#)
- [zxjdbc](#) for Jython

## Supported Versions and Features

SQLAlchemy supports MySQL starting with version 4.1 through modern releases. However, no heroic measures are taken to work around major missing SQL features - if your server version does not support sub-selects, for example, they won't work in SQLAlchemy either.

## SQLAlchemy 1.3 Documentation

Release: **1.3.0b2** **BETA RELEASE** | Release Date: January 25, 2019SQLAlchemy 1.3  
Documentation**BETA RELEASE**[Contents](#) | [Index](#)Search terms: 

ads via Carbon

## Dialects

[PostgreSQL](#)[MySQL](#)[SQLite](#)**Oracle**

- [Support for the Oracle database.](#)
- [Connect Arguments](#)
- [Auto Increment Behavior](#)
- [Identifier Casing](#)
- [LIMIT/OFFSET Support](#)
- [RETURNING Support](#)
- [ON UPDATE CASCADE](#)
- [Oracle 8 Compatibility](#)
- [Synonym/DBLINK Reflection](#)

## Oracle

Support for the Oracle database.

## DBAPI Support

The following dialect/DBAPI options are available. Please refer to individual DBAPI sections for connect information.

- [cx-Oracle](#)
- [zxJDBC for Jython](#)

## Connect Arguments

The dialect supports several `create_engine()` arguments which affect the behavior of the dialect regardless of driver in use.

- `use_ansi` - Use ANSI JOIN constructs (see the section on Oracle 8). Defaults to `True`. If `False`, Oracle-8 compatible constructs are used for joins.
- `optimize_limits` - defaults to `False`, see the section on LIMIT/OFFSET.
- `use_binds_for_limits` - defaults to `True`, see the section on LIMIT/OFFSET.

## SQLAlchemy 1.3 Documentation

Release: **1.3.0b2** **BETA RELEASE** | Release Date: January 25, 2019SQLAlchemy 1.3  
Documentation**BETA RELEASE**[Contents](#) | [Index](#)Search terms: 

ads via Carbon

## Dialects

[PostgreSQL](#)[MySQL](#)[SQLite](#)[Oracle](#)**Microsoft SQL Server**

- Support for the Microsoft SQL Server database.
- Auto Increment Behavior / IDENTITY Columns
  - Controlling "Start" and "Increment"
  - INSERT behavior
- MAX on VARCHAR / NVARCHAR
- Collation Support

## Microsoft SQL Server

Support for the Microsoft SQL Server database.

## DBAPI Support


The following dialect/DBAPI options are available. Please refer to individual DBAPI sections for connect information.

- [PyODBC](#)
- [mxODBC](#)
- [pymssql](#)
- [zxJDBC for Jython](#)
- [adodbapi](#)

## Auto Increment Behavior / IDENTITY Columns

SQL Server provides so-called "auto incrementing" behavior using the `IDENTITY` construct, which can be placed on any single integer column in a table. SQLAlchemy considers `IDENTITY` within its default "autoincrement" behavior for an integer primary key column, described at `Column.autoincrement`. This means that by default, the first integer primary key column in a `Table` will be considered to be the identity column and will generate DDL as such:

# Connecting to Databases with SQLAlchemy



**Connector**



**Connection String**

# Anatomy of a Connection String

```
engine = create_engine('sqlite:///sqlalchemy_sqlite.db')
```



# Anatomy of a Connection String

```
engine = create_engine('sqlite:///sqlalchemy_sqlite.db')
```



# Anatomy of a Connection String

```
sqlite:///sqlalchemy_sqlite.db
```



# Anatomy of a Connection String

`dialect:///`





# Anatomy of a Connection String

`dialect:///dbname`



# Anatomy of a Connection String

`dialect[+driver]://user:password@hostname/dbname`



# Anatomy of a Connection String

`dialect[+driver]://user:password@hostname/dbname[?key=value]`



```
engine_sqlite = create_engine('sqlite:///importing_sqlite.db')  
  
engine_postgres = create_engine('postgresql://xavier:postgres@localhost:5432/importing_postgres')  
  
engine_mysql = create_engine('mysql+mysqlconnector://root:mysql@localhost:3306/importing_mysql')
```



## A SQL Query

```
SELECT OwnerUserId,  
       SUM(AnswerCount) AS 'TotalAnswers',  
       SUM(ViewCount) AS 'TotalViews'  
FROM posts  
WHERE owneruserid is not NULL  
GROUP BY OwnerUserId  
ORDER BY 'TotalAnswers' DESC  
LIMIT 10;
```



## A SQL Query

```
SELECT OwnerUserId,  
       SUM(AnswerCount) AS 'TotalAnswers',  
       SUM(ViewCount) AS 'TotalViews'  
FROM posts  
WHERE owneruserid is not NULL  
GROUP BY OwnerUserId  
ORDER BY 'TotalAnswers' DESC  
LIMIT 10;
```



```
SELECT OwnerUserId,  
       SUM(AnswerCount) AS 'TotalAnswers',  
       SUM(ViewCount) AS 'TotalViews'  
FROM posts  
WHERE owneruserid is not NULL  
GROUP BY OwnerUserId  
ORDER BY 'TotalAnswers' DESC  
LIMIT 10;
```



```
mysql> SELECT OwnerUserId,  
-> SUM(AnswerCount) AS 'TotalAnswers',  
-> SUM(ViewCount) AS 'TotalViews'  
-> FROM posts  
-> WHERE owneruserid is not NULL  
-> GROUP BY OwnerUserId  
-> ORDER BY 'TotalAnswers' DESC  
-> LIMIT 10;
```





```
mysql> SELECT OwnerUserId,  
-> SUM(AnswerCount) AS 'TotalAnswers',  
-> SUM(ViewCount) AS 'TotalViews'  
-> FROM posts  
-> WHERE owneruserid is not NULL  
-> GROUP BY OwnerUserId  
-> ORDER BY 'TotalAnswers' DESC  
-> LIMIT 10;
```



```
mysql> SELECT OwnerUserId,  
-> SUM(AnswerCount) AS 'TotalAnswers',  
-> SUM(ViewCount) AS 'TotalViews'  
-> FROM posts  
-> WHERE owneruserid is not NULL  
-> GROUP BY OwnerUserId  
-> ORDER BY 'TotalAnswers' DESC  
-> LIMIT 10;
```

OwnerUserId	TotalAnswers	TotalViews
5	1	448
36	5	528
51	NULL	NULL
22	NULL	NULL
66	4	1335
64	0	543
63	2	322
-1	NULL	NULL
84	42	37229
96	5	311

```
result = engine.execute("""SELECT OwnerUserId,  
    SUM(AnswerCount) AS 'TotalAnswers',  
    SUM(ViewCount) AS 'TotalViews'  
FROM posts  
WHERE owneruserid is not NULL  
GROUP BY OwnerUserId  
ORDER BY 'TotalAnswers' DESC  
LIMIT 10;""").fetchall()  
  
pd.DataFrame(result, columns=[ 'OwnerUserId', 'TotalAnswers',  
    'TotalViews' ])
```

## Executing a SQL Query

Use **execute** with a SQL statement

- Error prone

Use SQLAlchemy ORM




# Object-relational Mapping

Programming technique for converting data between incompatible type systems in object-oriented programming languages



# Object-relational Mapper



## Table

- Represents a table in the database

## Mapper

- Maps a Python class to a table

## Class

- Object that defines how a record maps to an object

## Classical Mapping or Declarative API

```
from sqlalchemy import Table, MetaData, Column, Integer, String
from sqlalchemy.orm import mapper

metadata = MetaData()
```

## Classical Mapping

Mapped class using the **mapper** function

- Original class mapping

Create **MetaData**

- What the database looks like for SQLAlchemy



```
tags = Table('Tags', metadata,
    Column('Id', Integer, primary_key=True),
    Column('Count', Integer),
    Column('ExcerptPostId', Integer),
    Column('TagName', String(255)),
    Column('WikiPostId', Integer))
```

## Classical Mapping

### Define our **Table**

- Specify each **Column**, with the type
- Other attributes

```
class Tags(object):  
    def __init__(self, Count, ExcerptPostId, TagName, WikiPostId):  
        self.Count = Count  
        self.ExcerptPostId = ExcerptPostId  
        self.TagName = TagName  
        self.WikiPostId = WikiPostId  
  
tags_mapper = mapper(Tags, tags)
```

## Classical Mapping

Define **Tags** class

- Initialize

Associate via mapper **function**





```
larger_tags = tags.select(Tags.Count > 1000)
larger_tags
engine.execute(larger_tags).fetchall()
```

## Classical Mapping

### Create a query

- Use **select** and a condition

### Execute



# Declarative

## ORM

### Typically used system

- Provided by SQLAlchemy ORM

### Define classes

### Mapped to relational database tables

### Series of extensions

- On top of the mapper construct



```
from sqlalchemy.orm import sessionmaker  
session = sessionmaker()  
session.configure(bind=engine)  
my_session = session()
```

# Session

## Core concepts of SQLAlchemy

### Establishes and maintains conversations

- Between our program and the database
- Entry point for queries



```
my_session.query(Tags).all()  
len(my_session.query(Tags).all())  
my_session.query(Tags).first()  
my_session.query(Tags).first().TagName  
my_session.query(Tags.Id, Tags.TagName).first()
```

## Session

Can use session to **query**

- List of **Tags** objects
- Nicer representation?



```
from sqlalchemy.ext.declarative import declarative_base  
Base = declarative_base()
```

## Declarative Base

### **Base class for declarative class definitions**

- Define models
- Connect them to the database



```
class Users(Base):  
    __tablename__ = 'users'  
  
    Id = db.Column(db.Integer, primary_key = True)  
    Reputation = db.Column(db.Integer)  
    CreationDate = db.Column(db.DateTime)  
    DisplayName = db.Column(db.String(255))  
    LastAccessDate = db.Column(db.DateTime)  
    ...
```

## Define a Model

### Class

- Inherit from `declarative_base`
- Specify `__tablename__`
- Columns, with types



# Define a Model

```
class Users(Base):
    __tablename__ = 'users'

    Id = db.Column(db.Integer, primary_key = True)
    Reputation = db.Column(db.Integer)
    CreationDate = db.Column(db.DateTime)
    DisplayName = db.Column(db.String)
    LastAccessDate = db.Column(db.DateTime)
    WebsiteUrl = db.Column(db.String)
    Location = db.Column(db.String)
    AboutMe = db.Column(db.String)
    Views = db.Column(db.Integer)
    UpVotes = db.Column(db.Integer)
    DownVotes = db.Column(db.Integer)
    AccountId = db.Column(db.Integer)

    def __repr__(self):
        return "<{0} Id: {1} - Name: {2}>".format(self.__class__.__name__,
self.Id, self.DisplayName)
```



```
my_session.query(Users).first()  
type(my_session.query(Users).first())  
my_session.query(Users).first().DisplayName  
for each_user in my_session.query(Users):  
    print(each_user)
```

## Querying

Can **query** using the **session**

Returns the **\_\_repr\_\_** of the object

Can access the fields

Iterate over the results





```
the_query = my_session.query(Users)
type(the_query)
print(the_query)
engine_echo =
db.create_engine('mysql+mysqlconnector://root:mysql@localhost:3306/sqlalchemy_mysql', echo=True)
connection_echo = engine_echo.connect()
session_echo = sessionmaker(bind=engine_echo)()
session_echo.query(Users).first()
```

## The Query and Echo Parameter

### What query is executed?

- Print the **query**

### Use the echo parameter when creating engine

- Execute and get the SQL statements



```
my_session.query(Users).filter_by(DisplayName='Community').all()
```

```
my_session.query(Users).filter(Users.DisplayName=='Community').all()
```

## Filtering Results

Possible to use **filter\_by**

- Can be prone to confusion

Use **filter**

- Explicit



```
my_session.query(Users.DisplayName).filter(Users.DisplayName.like(' %Comm%')).all()
```

```
my_session.query(Users.DisplayName).filter(Users.DisplayName.contains('Comm')).all()
```

## ClauseElements

**What we just passed to `filter`**

- Base class for elements of a programmatically constructed SQL expression

**For example `like`, but there are many others**



```
from sqlalchemy import func  
  
dir(func)  
  
tags_count =  
my_session.query(func.sum(Tags.Count)).scalar()
```

## Functions

Many functions available through **func**

- Function generator

Use **scalar** to return a single element



```
my_session.query(Users.DisplayName, db.cast((Users.UpVotes -  
Users.DownVotes), db.Numeric(12, 2)).label('vote_difference') ,  
Users.UpVotes, Users.DownVotes).all()
```

## Operators and Labels

### Perform operations

- Calculate a difference between two columns
- Use **cast**

Create a **label** to refer to the column



```
my_session.query(Users.DisplayName, db.cast((Users.UpVotes -  
Users.DownVotes), db.Numeric(12, 2)).label('vote_difference'),  
Users.UpVotes, Users.DownVotes).limit(5).all()
```

## Limiting Results

**Control how many results are retrieved**

- With **limit**



```
my_session.query(Users.DisplayName, db.cast((Users.UpVotes -
Users.DownVotes), db.Numeric(12, 2)).label('vote_difference'),
Users.UpVotes,
Users.DownVotes).order_by('vote_difference').limit(5).all()
```

```
my_session.query(Users.DisplayName, db.cast((Users.UpVotes -
Users.DownVotes), db.Numeric(12, 2)).label('vote_difference'),
Users.UpVotes,
Users.DownVotes).order_by(db.desc('vote_difference')).limit(5).all()
```

## Ordering Results

Use **order\_by** to sort results

- Ascending (default)
- Descending (**desc**)



```
my_session.query(Users.DisplayName).filter(Users.DisplayName ==  
'Community', Users.DownVotes.between(300,600)).all()
```

```
my_session.query(Users.DisplayName,  
Users.DownVotes).filter(db.or_(Users.DisplayName == 'Community',  
Users.DownVotes.between(300,600))).all()
```

## Conjunctions

### Filter by multiple statements

- Using **and\_**, **or\_**, and **not\_**
- And is the default, but better to add it





```
class Posts(Base):
    __tablename__ = 'posts'
    Id = db.Column(db.Integer(), primary_key=True)
    Title = db.Column(db.String(255), nullable=False)
    ViewCount = db.Column(db.Integer(), default=1000)

    PostTypeId = db.Column(db.Integer(), default=True)
    OwnerUserId = db.Column(db.Integer())
```

## Define Posts

The **Users** can have many **Posts**

Use it to **join** data



```
my_session.query(Users, Posts).filter(Users.Id ==  
Posts.OwnerUserId).limit(1).all()
```

## Implicit Join

Do not use the **join**

Returns both **Users** and **Posts**

- Where **Id** matches **OwnerUserId**



```
my_session.query(Users, Posts).join(Posts, Users.Id ==  
Posts.OwnerUserId).limit(1).all()
```


## Explicit Join

Use **join**

From **Users** to **Posts**



# Hierarchical Tables



## Tables that refer to themselves

### Posts

- An answer refers to a question
- Both are posts
- ParentId to Id

### Self join

- Using alias

```
class Posts(Base):
    __tablename__ = 'posts'
    Id = db.Column(db.Integer(), primary_key=True)
    Title = db.Column(db.String(255), nullable=False)
    ViewCount = db.Column(db.Integer(), default=1000)

    PostTypeId = db.Column(db.Integer(), default=True)
    OwnerUserId = db.Column(db.Integer())
    __table_args__ = {'extend_existing': True}
    AnswerCount = db.Column(db.Integer)
    ParentId = db.Column(db.Integer)
```

## Extending a Model

Extend a model using **\_\_table\_args\_\_**

- With **extend\_existing** set to **True**



```
my_session.query(Posts.Id, Posts.Title, Posts.AnswerCount).filter(Posts.Id == 14).all()

my_session.query(Posts.Id).filter(Posts.ParentId == 14).all()

from sqlalchemy.orm import aliased

Questions = aliased(Posts)

my_session.query(Posts.Id, Questions.Id, Posts.ViewCount, Posts.Score, Questions.Score).filter(Posts.Id == Questions.ParentId).order_by(db.desc(Posts.ViewCount)).limit(10).all()
```

## Hierarchical Tables

Create an alias with **aliased**

Use the aliased table

- To **join** data



# Takeaway



## Querying

- Common operation

## Picking a database

- Connector
- Connection string

# Takeaway



## Query with SQL statements

- With `engine.execute`

## Object-relational mapper

- Classical mappings
- Declarative API



# Takeaway



## Declarative API

- Session & declarative\_base
- Define a model
  - Base, \_\_tablename\_\_, and columns

## Query

- Filter, functions, sort, limit, join, self join